



Lecture 15 Quiz



7/7 points earned (100%)

Quiz passed!

[Continue Course](#)[Back to Week 15](#)

1 / 1
points

1. The objective function of an autoencoder is to reconstruct its input, i.e., it is trying to learn a function f , such that $f(x) = x$ for all points x in the dataset.

It makes sense to learn a mapping from x to some target t for solving a classification or regression problem, but why do we want to learn a mapping that takes x to x ? It seems like a silly thing to do!



While the objective is to reconstruct x , what we really care about is learning good representations in the hidden layers which come as a by-product of doing this.



Correct

This is the most important reason why we want to learn autoencoders.



We want to have a decoder network, so that it is easy to get a reconstruction of the data.



Due to the constraints in the network, x will never be mapped exactly to x but to something close-by, say \hat{x} . We can then use \hat{x} as a better representation of x .



While the objective is to learn f such that $f(x) = x$, we might end up learning a mapping to a totally different space in which $f(x)$ might be a much better representation of the input.



1 / 1
points

2. The process of autoencoding a vector seems to lose some information since the autoencoder cannot reconstruct the input exactly (as seen by the blurring of reconstructed images reconstructed from 256-bit codes). In other words, the intermediate representation appears to have less information than the input representation. In that case, why is this intermediate representation more useful than the input representation ?



The intermediate representation loses some information, but retains what is most important. We hope that this retained information is "semantic". The intermediate representation will then be a more direct way of representing semantic content.



Correct



The intermediate representation has more noise.



The intermediate representation actually has more information than the inputs.



The intermediate representation is more compressible than the input representation.



1 / 1
points

3. Why is it harder to train deep autoencoders compared to shallow ones ?



The objective function is intractable for deep autoencoders.



Starting with small initial weights, backpropagated gradients become very small going down through the network, so the bottom layers do not get much gradient.



Correct

Due to the presence of saturating activation functions, the gradients may start decreasing rapidly going down the network.



Backpropagation is not exact for deep autoencoders.



The Markov chain involved in estimate negative phase sufficient statistics does not mix well.



1 / 1
points

4. In all the autoencoders discussed in the lecture, the decoder network has the same number of layers and hidden units as the encoder network, but arranged in reverse order. Brian feels that this is not a strict requirement for building an autoencoder. He insists that we can build an autoencoder which has a very different decoder network than the encoder network. Which of the following statements is correct?



Brian is correct, as long as the decoder network has **at most** as many parameters as the encoder network.



Brian is correct. We can indeed have any decoder network, as long as it produces output of the same shape as the data, so that we can compare the output to the original data and tell the network where it's making mistakes.



Correct

An autoencoder is just a neural net trying to reconstruct its input. There is hardly any restriction on the kind of encoder or decoder to be used.



Brian is correct, as long as the decoder network has the same number of parameters as the encoder network.



Brian is mistaken. The decoder network must have the same architecture. Otherwise backpropagation will not work.



1 / 1
points

5. Another way of extracting short codes for images is to hash them using standard hash functions. These functions are very fast to compute, require no training and transform inputs into fixed length representations. Why is it more useful to learn an autoencoder to do this ?



Autoencoders can be used to do **semantic** hashing, where as standard hash functions do not respect semantics , i.e, two inputs that are close in meaning might be very far in the hashed space.



Correct

- ☐ Autoencoders have smooth objective functions whereas standard hash functions have no concept of an objective function.
 - ☐ For an autoencoder, it is possible to invert the mapping from the hashed value to the reconstruct the original input using the decoder, while this is not true for most hash functions.
 - ☐ Autoencoders have several hidden units, unlike hash functions.
-



1 / 1
points

6. RBMs and single-hidden layer autoencoders can both be seen as different ways of extracting one layer of hidden variables from the inputs. In what sense are they different ?



RBMs define a probability distribution over the hidden variables conditioned on the visible units while autoencoders define a deterministic mapping from inputs to hidden variables.



Correct

RBMs define a joint density $P(v, h)$. Given v , we can compute $P(h|v)$. Autoencoders define a function $h = f(v)$.



The objective function and its gradients are intractable to compute exactly for RBMs but can be computed efficiently exactly for autoencoders.



Correct

The objective function for RBMs is log-likelihood. This is intractable to compute due to the partition function. Its gradients are hard to compute for similar reasons and computing them approximately requires CD or other MCMC methods. Autoencoders usually have tractable objectives such as squared loss or cross entropy which are easy to compute and differentiate.



RBMs work only with binary inputs but autoencoders work with all kinds of inputs.



Un-selected is correct



RBMs are undirected graphical models, but autoencoders are feed-forward neural nets.



Correct



1 / 1
points

7. Autoencoders seem like a very powerful and flexible way of learning hidden representations. You just need to get lots of data and ask the neural network to reconstruct it. Gradients and objective functions can be exactly computed. Any kind of data can be plugged in. What might be a limitation of these models ?



The hidden representations are noisy.



The inference process for finding states of hidden units given the input is intractable for autoencoders.



There is no simple way to incorporate uncertainty in the hidden representation $h = f(v)$. A probabilistic model might be able to express uncertainty better since it is being made to learn $P(h|v)$.



Correct



Autoencoders cannot work with discrete-valued inputs.



